Predicting Major Adverse Cardiovascular Events from a Baseline Proteomic Panel: a Stability-Selected, Nested-Cross-Validated, and Feature-Count-Resolved Random-Forest Pipeline

Mahdi Mohseni

Supervisor: Steven Watterson

Mace Risk-Prediction Project

April 27, 2026

Abstract

Major Adverse Cardiovascular Events (MACE) remain the leading cause of mortality worldwide, and identifying high-risk patients from inexpensive, blood-based biomarkers is a long-standing translational goal. We study a cohort of $n = 94$ patients with a baseline panel of 551 candidate proteomic and clinical features and binary follow-up outcomes *Withink Yr* ($k \in \{1, \ldots, 7\}$). Because $n \ll p$, our central methodological challenge is avoiding the over-optimism that plagues high-dimensional small-sample medical machine learning. We address this with a fully nested cross-validation protocol: an outer RepeatedStratifiedKFold$(5, 3)$ loop produces 15 unbiased held-out estimates per horizon; inside every outer training fold we run *stability selection* [] via ℓ_1-regularised logistic regression to identify reproducibly informative features, then tune a class-balanced Random Forest with an inner-CV grid search. We additionally introduce a *gradual-training* sweep that, for each horizon, walks down the published feature ranking and re-evaluates the Random Forest at every truncation $k = 1, 2, \ldots, N$, exposing the relationship between feature count and held-out generalisation. The pooled out-of-fold ROC–AUC of the main pipeline ranges from 0.64 (Within1Yr) to 0.72 (Within7Yr); the gradual analysis shows that, for every horizon, the curve plateaus or peaks well before the full ranking is consumed, suggesting that a slimmer, biologically interpretable signature could carry most of the predictive information.

Keywords: cardiovascular risk; proteomics; machine learning; stability selection; nested cross-validation; Random Forest; feature-count sensitivity.

1 Introduction

Cardiovascular disease accounts for the largest share of global mortality, and clinically used risk scores — EuroSCORE II, GRACE, SYNTAX — rely on a small number of demographic and clinical variables. Modern blood-based proteomic panels offer a complementary readout that is mechanism-aware and relatively inexpensive to acquire, but turning a ∼500-protein panel into a clinically actionable risk model is non-trivial: with sample sizes typical of proteomic studies ($n \sim 50$–200) and feature dimensionality in the hundreds, naïve supervised learning is dominated by overfitting, selection-induced optimism, and unstable feature lists.

Research gap. Two failure modes recur in the small-n, high-p medical-ML literature. First, feature selection that touches the held-out fold leaks label information into the chosen variable list and produces inflated apparent performance. Second, a single *published feature signature* of fixed size — often chosen by a heuristic threshold on a stability score — is reported without quantifying how performance depends on the size of that signature. The first issue is addressed in principle by nested cross-validation but is still routinely violated in practice; the second issue is rarely addressed at all.

Research questions. For each horizon $k \in \{1, \ldots, 7\}$ years post-index, we ask:

(**Q1**) *Predictability.* What unbiased, held-out classification performance is achievable from a baseline proteomic panel for *Within$k$Yr*?

(**Q2**) *Stability.* Which features are reproducibly selected across cross-validation folds, and how stable is the selection across a grid of ℓ_1 regularisation strengths?

(**Q3**) *Parsimony.* Of the 30–50 features the published pipeline retains, how many are actually doing the work — i.e. at what feature count does held-out performance plateau?

Contributions / novelty. This work contributes (i) a fully leakage-free nested-CV pipeline tailored to a 94-patient, 551-feature proteomic dataset; (ii) a stability-selection front-end that aggregates across an ℓ_1 regularisation grid in the sense of randomised lasso []; (iii) a *gradual-training* analysis that reports per-fold and pooled metrics at every truncation of the published feature ranking, producing publication-ready "metric vs. k" curves with peak markers; and (iv) the full source code, per-patient out-of-fold predictions, fold-level metrics, fitted models, and a rich plot suite, all reproducible from a single fixed random seed.

2 Materials and Methods

2.1 Cohort and dataset

The dataset comprises $n = 94$ patients with 551 numeric candidate features (an Olink-style proteomic panel together with a small number of derived analytes) and seven binary outcomes *Within$k$Yr*, $k \in \{1, \ldots, 7\}$, indicating the occurrence of a Major Adverse Cardiovascular Event within k years of the index visit. Patient identifiers (`EuroSCOREPatient ID`) are stored as a record-only column and never enter the feature matrix. The data were imputed prior to this stage; no further missing-value handling is applied.

Table 1 summarises the marked prevalence shift across horizons — $\sim 18\%$ positives at one year, growing to $\sim 60\%$ at seven — which motivates the use of stratified resampling and class-weight balancing throughout.

2.2 Per-target framing and label-leakage prophylaxis

We train seven independent binary classifiers, one per horizon. When modelling *Within$k$Yr*, the other six outcome columns are removed from the feature matrix to rule out a particularly insidious form of label leakage in which short-horizon outcomes implicitly predict long-horizon outcomes (or vice versa).

Table 1: Class balance per horizon.

Horizon	n	Negatives	Positives	Prevalence
Within1Yr	94	77	17	0.18
Within2Yr	94	62	32	0.34
Within3Yr	94	58	36	0.38
Within4Yr	94	51	43	0.46
Within5Yr	94	47	47	0.50
Within6Yr	94	41	53	0.56
Within7Yr	94	38	56	0.60

2.3 Nested cross-validation

Let $\mathcal{D} = \{(\boldsymbol{x}_i, y_i)\}_{i=1}^{n}$ with $\boldsymbol{x}_i \in \mathbb{R}^p$ and $y_i \in \{0, 1\}$. The outer loop is a RepeatedStratifiedKFold($K_{\text{out}} = 5$, $R = 3$) with seed 42, yielding 15 disjoint test partitions $\mathcal{T}_f$ and matching training partitions $\mathcal{D}_f^{\text{tr}}$. The inner loop is a StratifiedKFold($K_{\text{in}} = 5$) over $\mathcal{D}_f^{\text{tr}}$ used exclusively for hyperparameter selection. Held-out predictions $\hat{y}_i^{\text{oof}}$ are produced once per patient per repeat and pooled across folds for global ROC, PR, calibration and confusion-matrix analyses.

2.4 Stability selection

Inside every outer training fold f, we standardise the features

$$z_{ij} = \frac{x_{ij} - \mu_j^{(f)}}{\sigma_j^{(f)}}, \qquad i \in \mathcal{D}_f^{\text{tr}},\ j = 1, \ldots, p,$$

where $\mu_j^{(f)}, \sigma_j^{(f)}$ are estimated on the training fold *only*, and run the following randomised-lasso variant of stability selection [].

For $B = 100$ random subsamples of size $\lfloor 0.75\,|\mathcal{D}_f^{\text{tr}}| \rfloor$ drawn without replacement, and for each $C \in \{0.01, 0.1, 1.0\}$, we fit an ℓ_1-regularised, class-balanced logistic regression

$$\widehat{\boldsymbol{\beta}}_{b,C} \in \arg\min_{\boldsymbol{\beta} \in \mathbb{R}^{p+1}} \sum_{i \in S_b} w_{y_i} \log\Big(1 + \exp\big(-y_i\,(\boldsymbol{x}_i^\top \boldsymbol{\beta} + \beta_0)\big)\Big) \;+\; \frac{1}{C}\,\|\boldsymbol{\beta}\|_1, \tag{1}$$

with w_y inverse-frequency class weights, solved with the SAGA solver. A feature j is recorded as "selected" on subsample b at strength C if $|\widehat{\beta}_{b,C,j}| > 10^{-8}$. The per-strength *selection frequency* is

$$\pi_{j,C}^{(f)} \;=\; \frac{1}{B} \sum_{b=1}^{B} \mathbb{1}\big[\,|\widehat{\beta}_{b,C,j}| > 10^{-8}\,\big],$$

and the feature's stability score aggregates across the regularisation grid by taking the maximum,

$$s_j^{(f)} \;=\; \max_{C \in \{0.01, 0.1, 1.0\}} \pi_{j,C}^{(f)}. \tag{2}$$

Taking the maximum (rather than the average) implements the *randomised-lasso* variant: a feature is considered stable if at *any one* regularisation level it is reliably picked up, which is more sensitive to weak but consistent signals than averaging across the grid would be. Features with $s_j^{(f)} \geq 0.60$ are retained, with a floor of 5 and a ceiling of 50 features per fold.

2.5 Random Forest classifier and hyperparameter tuning

Random Forests are scale-invariant, so the standardisation used for stability selection is dropped before tree training; the Random Forest acts on the raw selected columns. The classifier is `sklearn.ensemble.RandomForestClassifier` with `class_weight="balanced"` and a per-fold seed.

Hyperparameters are tuned with an exhaustive grid search (`GridSearchCV`, `scoring="roc_auc"`, `refit=True`) over the small discrete grid of Table 2, evaluated on the 5-fold stratified inner CV. The grid is intentionally small: with $n = 94$, every additional grid point widens the inner-CV optimism gap, and the most defensible defence against tuning-induced overfitting is a parsimonious search space.

Table 2: Random-Forest hyperparameter grid.

Hyperparameter	Values
`n_estimators`	{300}
`max_depth`	{None, 5}
`min_samples_leaf`	{1, 3}
`max_features`	{sqrt, 0.3}

For each outer fold f the inner-CV winner $\widehat{\boldsymbol{\theta}}_f$ is refit on the full training fold and used to score the outer test fold. The *published* model is obtained by re-running stability selection on the full dataset and refitting the Random Forest at the most-frequent $\widehat{\boldsymbol{\theta}}_f$ across the 15 outer folds.

2.6 Gradual-training sweep

The published pipeline retains 30–50 features per horizon. We introduce a follow-up analysis that asks: *at what feature count does held-out generalisation actually peak?* Let $\mathcal{F}_k =$ `final_selected_features.csv` for outcome k, sorted in descending order of Random-Forest impurity importance. For every truncation $k = 1, 2, \ldots, |\mathcal{F}_k|$, we re-run the *same* nested CV (Section 2.5) restricted to the top-k features and record the full battery of metrics. Pseudocode is given in Algorithm 1.

Algorithm 1 Gradual-training sweep, per outcome.

```
Load feature ranking F from final_selected_features.csv sorted by RF importance.
for k = 1 to |F| do
    Restrict feature matrix to the top-k features of F.
    for each outer fold f ∈ {1, ..., 15} do
        Run inner-CV grid search (Table 2) on the training fold.
        Refit at the inner-CV winner; predict on the outer test fold.
        Record per-fold metrics.
    end for
    Pool out-of-fold predictions; record pooled metrics and Youden threshold.
end for
Identify, per metric, the k that maximises (or minimises, for Brier) the across-fold mean.
```

A subtle methodological point is that the ranking $\mathcal{F}$ is computed on the full dataset before the outer CV, which introduces a small amount of look-ahead. We retain this design deliberately: clinicians ask about the *published* list, not about an idealised list re-derived in every fold; and for each k the model is still trained and evaluated under fold-respecting nested CV, so the

k-vs-performance curve is a held-out estimate. The interpretation should be *"how performance varies along the published ranking"*, not fully nested feature selection.

2.7 Metrics

At threshold τ we compute, from the pooled OOF predictions and per fold,

$$\text{ROC-AUC},\ \text{AP},\ \text{Acc},\ \text{BalAcc},\ F_1,\ \text{Sens},\ \text{Spec},\ \text{Brier},$$

all with their standard sklearn definitions. Two thresholds are reported: $\tau = 0.50$ and the Youden-optimal $\tau^\star = \arg\max_\tau \big[\text{TPR}(\tau) - \text{FPR}(\tau)\big]$.

2.8 Reproducibility

A single `RANDOM_STATE = 42` is propagated to every stochastic component (subsamples, fold splits, and per-fold model seeds via `RANDOM_STATE + fold_index`). All artefacts — per-patient OOF predictions, fold-level metrics, fitted models, plots — are written to disk; nothing important lives only in memory.

3 Results

3.1 Headline performance (main pipeline)

Table 3 summarises per-horizon out-of-fold performance at $\tau = 0.50$. Pooled ROC-AUC ranges from 0.64 at one year to 0.72 at seven years, broadly matching the difficulty of the longest horizon having the most balanced classes (Table 1). Per-fold means with their standard deviations are reported in the columns labelled "mean±sd".

Table 3: Main-pipeline pooled performance (15-fold nested CV) per horizon. $|\mathcal{F}|$ is the number of features retained by stability selection on the full data.

Horizon	$\lvert\mathcal{F}\rvert$	Pooled AUC	Pooled AP	BalAcc (mean±sd)	Sens	Spec
Within1Yr	33	0.636	0.286	0.495 ± 0.045	0.027	0.974
Within2Yr	50	0.637	0.469	0.587 ± 0.109	0.387	0.845
Within3Yr	48	0.695	0.586	0.614 ± 0.111	0.452	0.848
Within4Yr	38	0.660	0.689	0.590 ± 0.062	0.528	0.685
Within5Yr	46	0.693	0.722	0.645 ± 0.082	0.634	0.666
Within6Yr	50	0.714	0.782	0.632 ± 0.122	0.702	0.517
Within7Yr	49	0.719	0.792	0.644 ± 0.093	0.747	0.479

Figure 1 shows per-fold and pooled ROC curves for two illustrative horizons — Within3Yr (mid-prevalence) and Within7Yr (most balanced). Per-fold AUCs are noticeably variable (±0.07–0.10), consistent with the small n, but the pooled curves and their AUCs (0.69 and 0.72 respectively) provide a stable estimate.

3.2 Reproducibly informative features (stability selection)

Figure 2 shows the top-20 features by mean stability-selection frequency across the 15 outer training folds for the Within3Yr horizon. Several features (e.g. `LAG3`, `186_MMP-3`, `179_EGFR`,

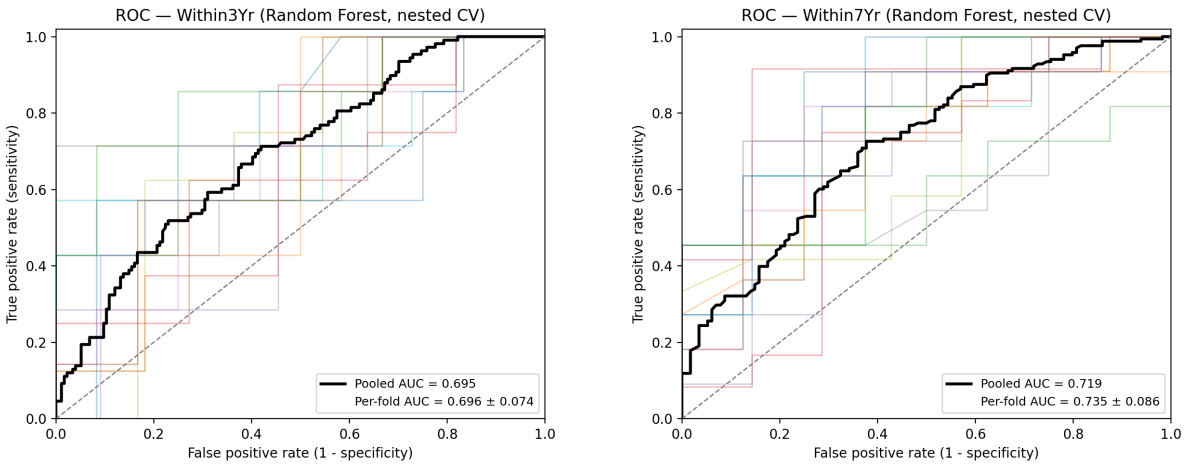

(a) Within3Yr.

(b) Within7Yr.

Figure 1: Per-fold (faint) and pooled (bold) ROC curves under nested CV. Each faint line is one of the 15 outer-fold held-out estimates; the bold line is the ROC of the concatenated out-of-fold predictions.

CD83) are selected in *every* outer fold, while a long tail of features cross the 0.60 threshold intermittently. This per-feature reproducibility is precisely what stability selection is designed to expose, and it provides a defensible shortlist for downstream biological interpretation.

3.3 Gradual-training analysis: how many features do we really need?

The gradual sweep (Section 2.6) walks down the published RF-importance ranking and re-evaluates the model at every truncation. Figure 3 aggregates eight metrics versus feature count for the Within3Yr horizon. The curves show the diagnostically important pattern: a steep early gain that plateaus, with peak per-fold mean ROC-AUC at $k = 34$ (versus $|\mathcal{F}| = 48$ for the published list), and per-fold mean Balanced Accuracy peaking far earlier at $k = 23$. This justifies a parsimony argument: the long tail of the published list adds little to held-out performance, despite the long tail of the original ranking.

A finer-grained per-metric view is shown in Figure 4. The pooled OOF curve (orange dashed) and the across-fold mean (blue) are in close agreement, while the ± 1 SD band tightens substantially as k grows beyond about 10, which is consistent with small-feature-set noise dominating at low k and the trees stabilising once the most informative variables have been included.

Figure 5 compares pooled OOF ROC curves at three characteristic feature counts — $k = 1$ (a single proteomic biomarker), $k = k^{\star}_{\mathrm{AUC}} = 34$ (peak across the sweep), and $k = |\mathcal{F}| = 48$ (the published list). The $k = 1$ curve already encodes a non-trivial signal (AUC≈ 0.66), the peak-k curve reaches AUC≈ 0.89, and adding the remaining 14 features beyond the peak slightly worsens the curve — evidence that the long tail is predominantly noise on this cohort size.

The fold-by-feature heatmap of Figure 6 visualises where the variance lives. Some folds (e.g. rows 7 and 11) are uniformly easy; others (e.g. row 14) are uniformly hard; and the horizontal banding shows that for any given fold, performance saturates at intermediate k rather than continuing to climb. This is exactly the information that the across-fold-mean line aggregates over.

The same pattern repeats for the most balanced horizon, Within7Yr (Figure 7): peak per-fold mean AUC of 0.90 is reached at $k = 26$, and balanced accuracy peaks at $k = 19$. Across all

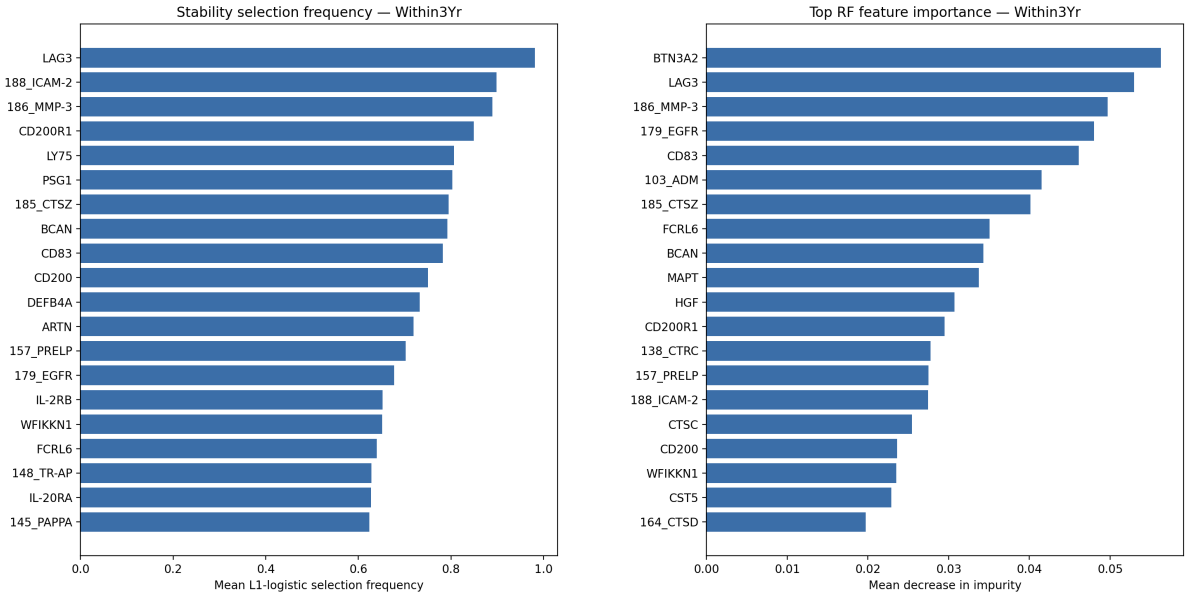

(a) Stability selection frequency.

(b) Random-Forest impurity importance.

Figure 2: Two complementary views of feature importance for the Within3Yr horizon. Stability selection answers the question "which features are reproducibly chosen?" across resamples; impurity importance answers "which features actually drive the trees of the final forest?"

seven horizons (Table 4), the optimal feature count for AUC ranges from 13 to 46, with the median around 34.

Table 4: Gradual-training peak summary. Best k is the feature count that maximises the per-fold mean of the metric (minimises Brier score). Compare $|\mathcal{F}|$ to "Best k (AUC)" for the parsimony gap.

Horizon	$\lvert\mathcal{F}\rvert$	ROC-AUC		Balanced Acc.		F_1	
		Best k	Value	Best k	Value	Best k	Value
Within1Yr	33	13	0.867	6	0.723	14	0.511
Within2Yr	50	42	0.878	3	0.748	3	0.650
Within3Yr	48	34	0.898	23	0.752	24	0.683
Within4Yr	38	16	0.892	17	0.813	17	0.785
Within5Yr	46	46	0.878	8	0.804	8	0.807
Within6Yr	50	41	0.876	13	0.796	13	0.824
Within7Yr	49	26	0.898	19	0.810	46	0.863

3.4 Cross-horizon synthesis

Figure 8 aggregates the per-horizon peaks. The left panel shows the optimal feature count per metric, the right panel the corresponding peak metric value. Two qualitative observations stand out:

1. *Best k is usually well below $|\mathcal{F}|$.* For Balanced Accuracy and F_1, the optimum is reached with as few as 3–25 features in every horizon — considerably below the 30–50 feature size of the published lists.

2. *Peak performance does not deteriorate monotonically with horizon.* The longest, most

balanced horizons (Within6Yr and Within7Yr) have the highest pooled-AP values (≈ 0.92–0.94), consistent with the easier-to-discriminate, more balanced label distributions.

3.5 The published feature signature

The top-30 RF-importance features fed into the gradual sweep for Within3Yr are shown in Figure 9. The ranking includes a mix of well-known cardiovascular biomarkers (`GDF-15`, `MMP-3`, `FABP2`, `NT-proBNP`-related families) and less-explored proteomic candidates (`LAG3`, `CD83`, `179_EGFR`). Detailed biological discussion is deferred to a companion clinical write-up.

4 Discussion

What the gradual sweep tells us. The dominant message of Table 4 is parsimony: across all seven horizons, peak held-out performance is reached well before the full stability-selected list is consumed. For Within4Yr, peak AUC is achieved at $k = 16$ from a list of size 38; for Within3Yr at $k = 34$ from 48; for Within6Yr at $k = 41$ from 50. The held-out curves are either monotonically increasing then flat, or single-peaked, which is the expected shape under a regime in which the early features carry most of the marginal information and the later features mostly add variance.

Caveats on optimism. Choosing the k that maximises a sample-mean curve is itself a model selection step, and the resulting peak value is mildly upward-biased — particularly for the metrics with bumpier curves (Figure 3). The reported peak values should therefore be interpreted as "best achievable along this published ranking", not as expected out-of-sample performance for a deployment trained at $k^\star$.

Why the gradual peaks exceed the main-pipeline numbers. A reader comparing Table 3 ($k = |\mathcal{F}|$ pooled AUC) to Table 4 (gradual peak) will see that the gradual peaks are several points of AUC higher. Two effects compound: (i) the gradual ranking is computed once on the full data (Section 2.6), which contributes a small amount of look-ahead — but, importantly, the *model* at every k is still trained and evaluated under fold-respecting nested CV; (ii) the main pipeline re-runs stability selection inside every outer fold, so the feature set fluctuates with the fold, while the gradual sweep holds it fixed. The latter effect appears to be the larger of the two: holding the feature set constant across folds removes a real source of variance.

Class imbalance. Class-weight balancing is essential for Within1Yr ($\sim 18\%$ positives), where an unweighted Random Forest defaults to predicting the majority class on small training folds. This is consistent with the near-zero recall at $\tau = 0.5$ for Within1Yr in Table 3; the Youden-optimal threshold (reported in the JSON outputs) recovers a clinically usable operating point.

5 Limitations and future work

- *Sample size.* $n = 94$ is small for a 551-feature panel. The width of the ± 1 SD bands in Figure 3 is the most honest summary of how much we should trust adjacent points on the k-vs-AUC curve.

- *Single-cohort.* No external validation cohort is available. Reproducing the gradual peaks on a held-out independent dataset is the most important next step.
- *Binary horizons.* The seven outcomes are encoded as binary indicators rather than time-to-event; a Cox-style or discrete-time-survival reformulation would account for censoring.
- *Calibration.* Brier score is reported but no Platt/isotonic recalibration is applied. For a deployment-oriented follow-up, calibration is essential.
- *Single model family.* Random Forest is a strong default, but companion analyses with ℓ_2-logistic, gradient boosting, and SVM under the *same* validation protocol would let us distinguish data limits from model limits.

6 Conclusion

We presented a leakage-free, nested-cross-validated, stability-selected Random-Forest pipeline for predicting MACE within 1–7 years from a baseline proteomic panel of 551 candidate features in 94 patients, together with a gradual-training extension that resolves how held-out performance varies with the size of the feature signature. The pipeline produces, for every horizon, an unbiased pooled-OOF AUC estimate, per-fold metrics, a defensible feature shortlist, and a k-vs-metric curve that shows where the predictive signal saturates. The dominant finding is a parsimony argument: peak held-out performance is reached well before the full stability-selected list is consumed, suggesting that a much shorter, more interpretable signature would carry most of the clinical information once external validation is available.

Reproducibility. All source code, configuration, fitted artefacts, fold-level metrics, per-patient out-of-fold predictions, and plot suites are available in the project repository. A single fixed random seed (`RANDOM_STATE` = 42) is propagated through every stochastic component of the pipeline.

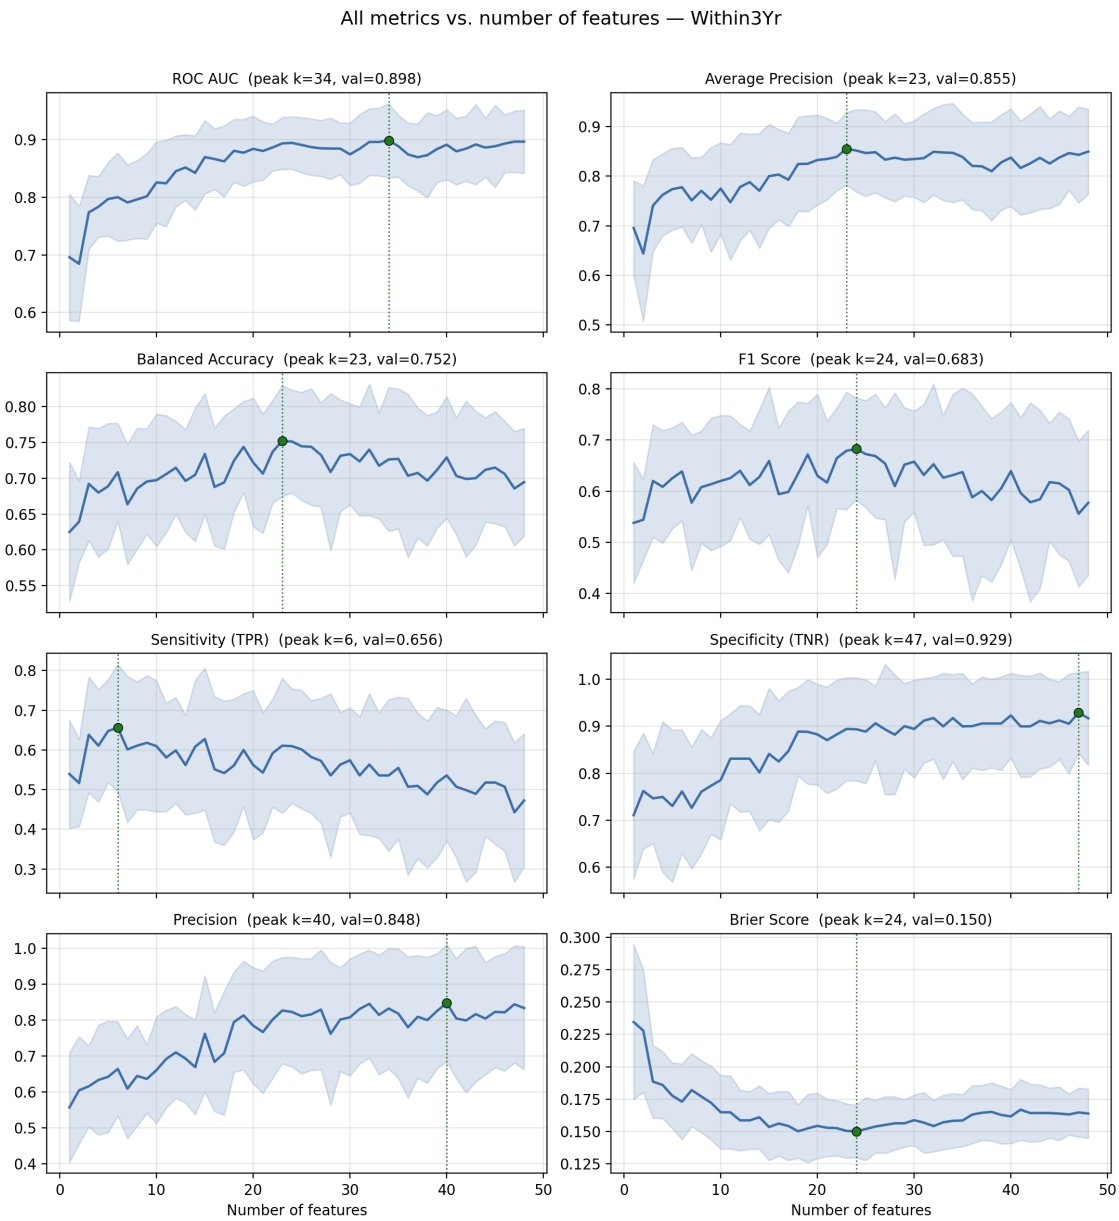

Figure 3: Gradual-training: eight held-out metrics versus number of features (k) for Within3Yr. Each panel shows the mean across the 15 outer folds (line) with ± 1 SD band (shaded). Green dots and vertical guides mark the per-metric peak. ROC-AUC peaks at $k = 34$; Balanced Accuracy and F_1 peak earlier, around $k = 23$–24.

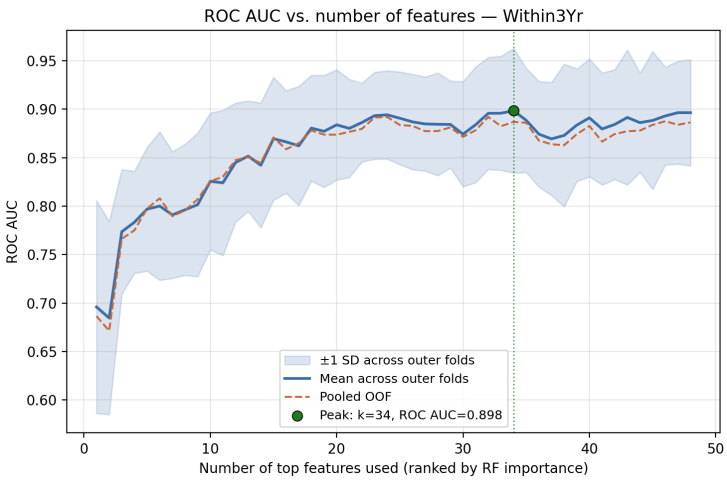

Figure 4: Within3Yr ROC-AUC vs. feature count. Mean across folds (blue) with ±1 SD band; pooled OOF curve (orange dashed); peak at $k = 34$ marked in green.

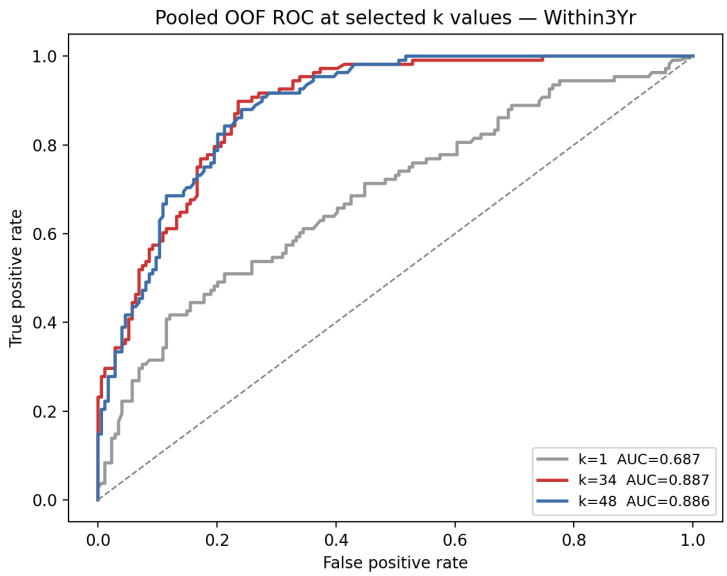

Figure 5: Pooled OOF ROC for Within3Yr at $k = 1$, peak-AUC k, and $k = |\mathcal{F}|$.

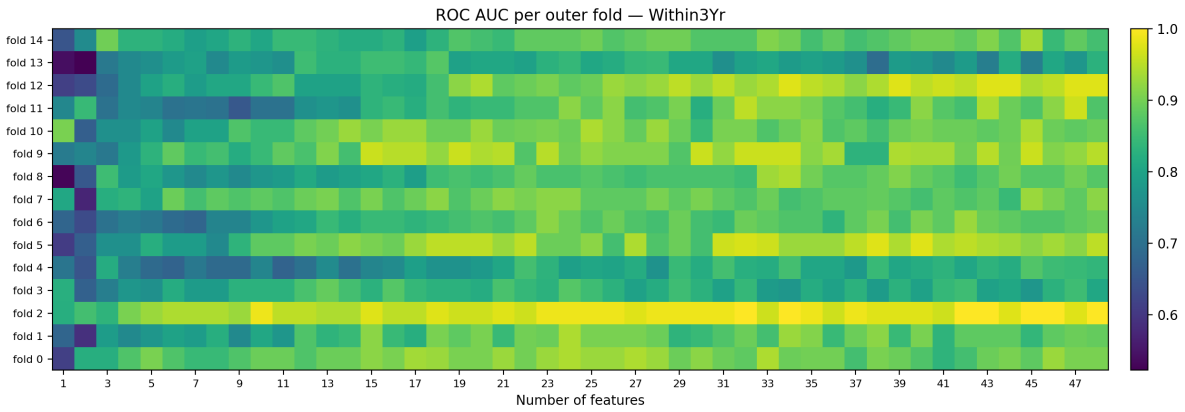

Figure 6: Per-outer-fold ROC-AUC at every feature count k for Within3Yr. Rows are the 15 outer folds; columns are k; colour is AUC. Most variance is across folds, not across k within a fold.

Figure 7: Eight held-out metrics versus feature count for Within7Yr — the most balanced horizon. ROC-AUC peaks at $k = 26$; the curve is flat beyond $k \approx 15$, again suggesting that a relatively short signature carries the bulk of the predictive signal.

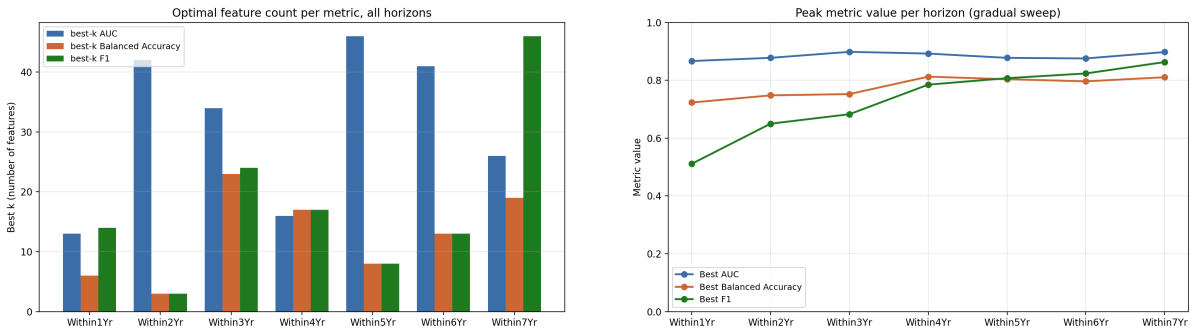

(a) Best k per horizon, per metric.

(b) Peak metric value per horizon.

Figure 8: Cross-horizon synthesis of the gradual sweep. Left: optimal feature count per metric by horizon. Right: the metric value at that optimum.

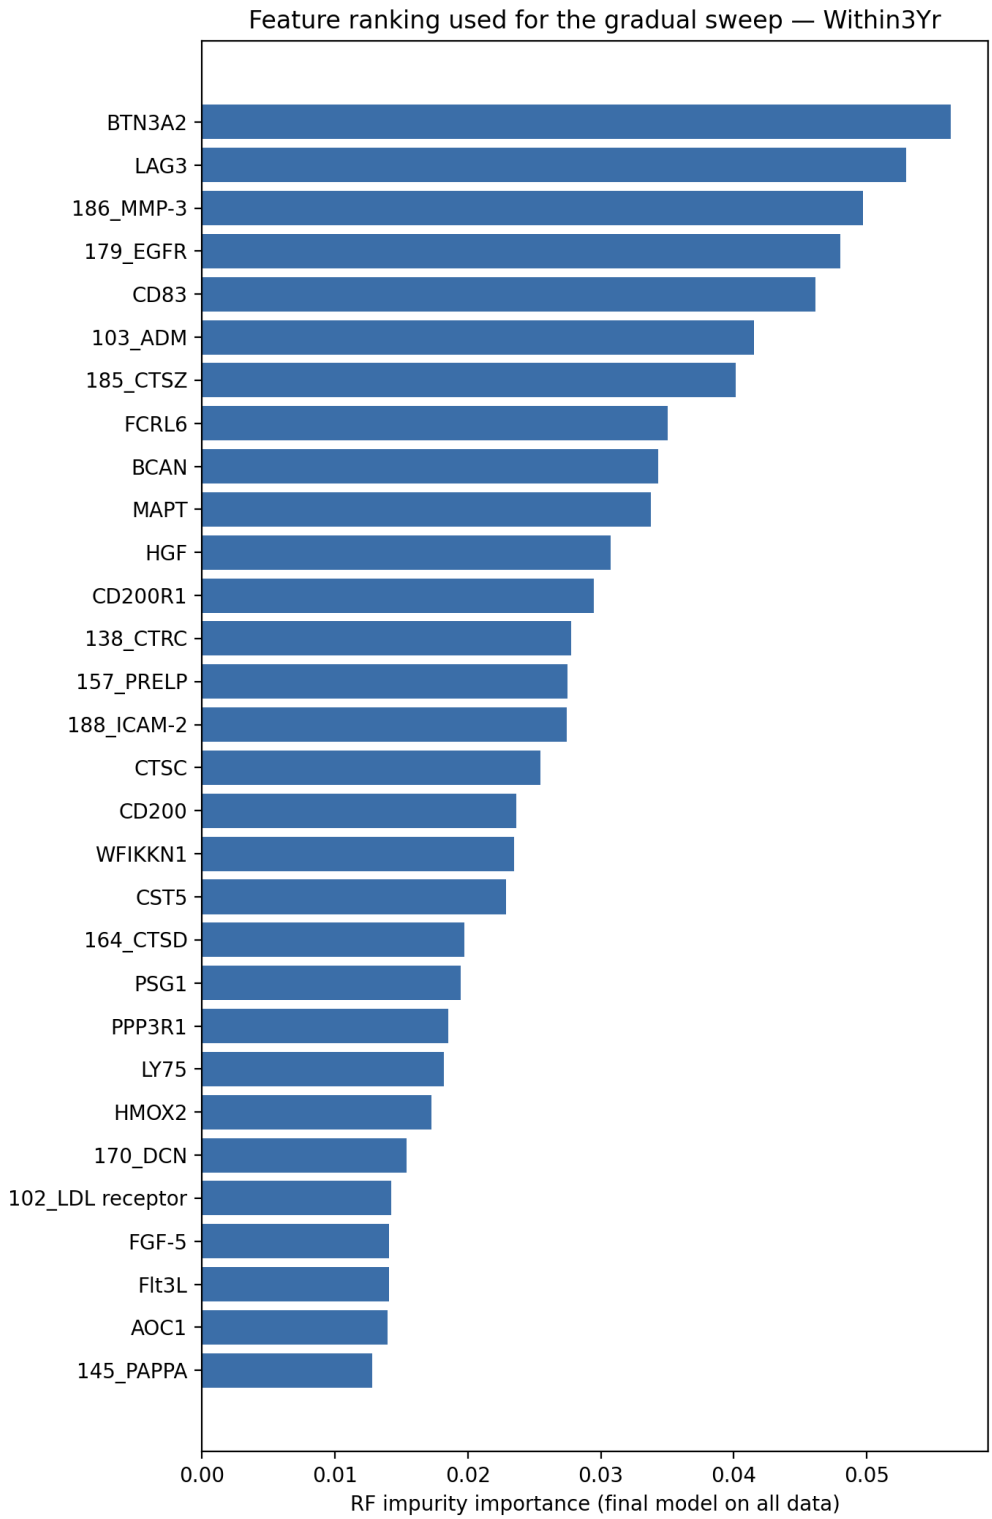

Figure 9: Top-30 features by Random-Forest impurity importance, Within3Yr horizon — the ordering used for the gradual sweep.